


The launch check for apps built with AI.

A complete briefing on what Walkthru is, how every part of it works, what each plan sells, where the business stands, and what a business co-founder would own.



Prepared for
A prospective co-founder, and the founder

State of the product
As of 24 September 2026, branch 1

Launch target
Tue 20 October 2026, 27 October at the latest

How to read this document

Every chapter is written twice over. The main text explains things in plain words, as if you have never built software. The grey boxes marked **Go deeper** give the technical detail for anyone who wants it. You can skip every grey box and still understand the whole business.

PLAIN VERSION

Read the main text only. It takes about 45 minutes and tells you what the product does, who pays for it, what it costs to run and what happens next.

GO DEEPER

Grey boxes name the actual files, numbers and rules in the code, for example `app/plans.py`, so an engineer can check every claim against the repository.

HONESTY NOTE

Nothing here is invented. Numbers come from the code, from measured runs, or from sourced research. Anything that is a guess is labelled **Estimate** with the arithmetic shown. Walkthru has **no paying customers yet**, so every revenue figure is a scenario, not a result.

Contents

01	Walkthru in one page	11	Is the free report real? An honest audit
02	The problem, and why now	12	Who benefits, and how
03	Glossary: every term explained	13	How it is built
04	Instant Scan	14	Market and competitors
05	The Chrome extension and AI test users	15	Business model and unit economics
06	The report, section by section	16	Initial investment and running costs
07	The services: SEO, GEO, security, email, prompts, badges	17	Go-to-market plan
08	Reading the numbers	18	What we build next
09	Safety: how the agent contains itself	19	Risks and open decisions
10	The plans, in full	20	The co-founder role
		A	Sources

Walkthru in one page

Walkthru checks whether a website is ready for real visitors before and after it launches. A founder pastes an address, and Walkthru answers three questions in one report, with a ranked list of fixes.

The question	How Walkthru answers it
1. Can a stranger get in?	An AI test user (a piece of software that behaves like a first-time visitor) tries a real task on the site, such as "sign up" or "find pricing", inside the owner's own Chrome browser. It thinks aloud, and the report shows exactly where it got stuck, with screenshots.
2. Can Google and AI search read you?	SEO checks (can Google find and list the pages) plus GEO checks (can ChatGPT, Claude, Perplexity and Google's AI answers read and quote the site).
3. Are you leaking anything?	Passive security checks: missing protective headers, weak cookies, exposed files, secret keys left in public code. Walkthru only looks; it never attacks.

On top of that: accessibility and speed checks, a signup email check, a single **Launch Ready score** from 0 to 100 with a public badge, and for paid plans a ready-made **fix prompt** the founder pastes into their AI coding tool (Cursor, Claude Code, Lovable, Bolt) to fix everything, then reruns Walkthru to prove it worked.

~11 s Instant Scan of a small site, measured today (10-page limit)	82% Seeded problems found on our hard test site (18 of 22)	171 Automated API tests passing, plus 30 extension tests	\$0 Model cost today: every call runs on free AI providers
--	--	--	--

Who it is for

- **Primary:** solo founders and two-person teams shipping apps they built with AI tools (Cursor, Lovable, Bolt, v0). No designer, no QA tester, often no idea whether a stranger can actually sign up.
- **Secondary:** small agencies that hand websites to clients and want proof the site works.

What makes it different

- **One report instead of four tools.** Competitors that send AI users do not check SEO, GEO or security. GEO tools never try a signup flow.
- **Logged-in pages without sharing a password.** The test runs in the owner's own browser, where they are already signed in. Every competitor we checked runs in its own cloud browser.
- **It does not lie.** Every finding must point at a real step, URL, header or snippet. When Walkthru stops itself for safety, it says so instead of blaming the site.

Where it stands (24 September 2026)

- Built and working locally: Instant Scan, extension, AI test users, reports, GEO, fix pack, fix prompt, rerun and compare, plan limits, data deletion, legal page drafts.
- Not done: payments (Dodo), deployment, domain, Chrome Web Store listing, the updated landing page. Plus-plan features (weekly watch, MCP server) come after launch.
- Prices set: Free, Launch Pack \$9 once, Pro \$19 a month (\$15 founding), Plus \$49 a month (\$39 founding).

The problem, and why now

Millions of people now build real web apps by describing them to an AI. They can ship in a weekend. What they cannot easily do is watch a stranger use it.

The pain, in the founder's words

- "I launched on Product Hunt, got 400 visitors and 3 signups. Was the signup broken, or is the product bad?"
- "My app is a React single-page app. Does Google even see my text? Does ChatGPT?"
- "Cursor wrote my auth. Did it leave an API key in the JavaScript?"

These are illustrative questions, not quotes from real users. Collecting real ones is one of the co-founder's first jobs (chapter 20).

How they deal with it today

Today's option	Why it falls short
Ask friends to try it	Friends are polite, already know the product, and only do it once.
Paid user testing (UserTesting, Maze)	Priced for companies. UserTesting sold for \$1.3B in 2022, proof the need is real, but far above an indie budget.
Free tools: Lighthouse, securityheaders.com, Google Search Console	Each covers one slice, speaks in jargon, and none tries the signup flow.
AI user-testing tools (Meerkat, CanaryUsers, Swarm)	Test journeys only. No SEO, GEO or security. Need credentials for logged-in pages.

Why now

- **AI app builders exploded.** Lovable passed \$400M in annual revenue in February 2026 and was near 8 million users by late 2025; Cursor passed \$2B in annual revenue in February 2026 (sources in the appendix). Estimates put the "vibe coding" tools market around \$4.7B in 2026. Every one of those apps needs a launch check.
- **Search is moving into AI answers.** People ask ChatGPT and Perplexity instead of Google. Many AI-built apps are JavaScript shells that AI crawlers see as empty pages. Almost nobody checks this.
- **Browser agents became cheap and reliable** in 2025 and 2026, so a "robot visitor" is now affordable to run on free AI models.

Glossary: every term explained

Read this once. The rest of the document uses these words freely.

Product words

AI persona / test user	A character the AI plays while using the site, such as "a busy first-time visitor who skims and gets impatient". Walkthru has four: first-time visitor , phone user (small screen, hates long forms), careful buyer (cares about pricing and trust), skeptic (a developer who reads error messages and distrusts vague copy). Different people get stuck in different places, so each persona finds different problems.
Journey / run	One test: one persona trying one goal on one site. A run is made of steps ; each step is one action like a click or typing into a field. Plans are counted in runs.
Goal	What the test user tries to do, typed by the owner: "Sign up for an account", "Find the pricing plans". Walkthru turns it into a short checklist before starting.
Think-aloud	The test user's running commentary at each step, like "Both buttons look the same, I will try the first one." It is what makes the report feel like watching a real person.
Instant Scan	The free, no-install check. Paste a URL, get a report in seconds. It reads public pages but does not click through a journey.
Extension	A small add-on installed in Chrome. It lets Walkthru act inside the owner's own browser tab.
Finding	One problem in the report, with a kind (UX, SEO, GEO, security, accessibility, performance), a severity (high, medium, low), evidence (where exactly) and a fix.
Fix pack	Copy-paste code blocks that fix GEO problems, built from the site's own pages: robots.txt rules, structured data, an llms.txt file, a rendering fix for the framework.
Agent fix prompt	One long instruction, generated from the report, that the owner pastes into their AI coding tool so it fixes every finding in order. Paid plans only.
Launch Ready score	One number from 0 to 100 combining UX, security, GEO, SEO and speed. 85 or more is "Launch ready", 60 to 84 "Almost ready", below 60 "Needs work".
Badge	A small image ("Walkthru, Launch Ready 92") the owner puts in their site footer or GitHub README. It links to the public report and updates on rerun. It is social proof for the owner and a backlink for us.
Rerun and compare	Run the same goal again after fixing. The report marks every earlier finding as fixed, still broken, new, or not re-checked.
Domain verification	Proof that the person testing a site owns it: a meta tag in the homepage or a file at <code>/ .well-known/walkthru.txt</code> . It unlocks deeper security checks and real form sends.
Safe stop	The run ended because the test user reached a button Walkthru will not press (pay, delete, send). This is Walkthru's rule, not the site's fault, and the report says so.

Search words

SEO	Search Engine Optimisation. Making pages easy for Google and Bing to find, understand and list: page titles, descriptions, headings, image descriptions, a sitemap, no accidental "do not index" rules.
GEO	Generative Engine Optimisation, also called AI search readiness. Making pages easy for AI assistants (ChatGPT, Claude, Perplexity, Google's AI Overviews) to fetch, understand and quote. Walkthru measures readiness; it never promises the site will be cited.
AI crawler / bot	The program an AI company sends to read web pages, such as OAI-SearchBot (ChatGPT search), ClaudeBot, PerplexityBot. Training bots (GPTBot, Google-Extended) are different: blocking them is a legitimate choice.
robots.txt	A text file at a site's root that tells crawlers what they may read. One wrong line can hide a whole site from Google and AI search.
llms.txt	A newer, optional file that summarises a site for AI models. Evidence of impact is weak, so Walkthru weights it only 5 of 100 points and labels it low impact.
Structured data / JSON-LD	A hidden block of machine-readable facts in a page ("this is a software product, its name is X, it costs Y"). Search engines and AI assistants use it to understand the page.
Content before JavaScript	Many AI-built sites send an empty page and draw the text with JavaScript afterwards. Most AI crawlers do not run JavaScript, so they see nothing. Walkthru shows exactly what they see.

Security, speed and accessibility words

Passive security	Only reading what any visitor already receives (headers, cookies, public files). No hacking attempts, no load tests. Legal and safe to run.
HTTPS / TLS	The padlock: an encrypted connection. A site on plain http sends passwords in the clear.
Security headers	Instructions a server sends with each page, like Content-Security-Policy (which scripts may run), X-Frame-Options (stops the page being embedded by a fake site).
Cookie flags	HttpOnly (JavaScript cannot read the login cookie), Secure (only sent over HTTPS), SameSite (not sent from other sites). Missing flags make session theft easier.
Leaked keys	Secret API keys accidentally shipped inside public JavaScript. Anyone can copy them and run up the owner's bill.
SPF, DMARC, MX	Email settings for a domain. Without them, signup and password-reset emails land in spam or never arrive.
Accessibility (a11y)	Whether people using screen readers or keyboards can use the site: labelled form fields, image descriptions, heading order. Checked with axe, the industry-standard tool, against WCAG A and AA rules.
Core Web Vitals	Google's speed and stability measures. LCP (Largest Contentful Paint): how long until the main content shows; good is under 2.5 s. CLS (Cumulative Layout Shift): how much the page jumps around; good is under 0.1. INP (Interaction to Next Paint): how fast it reacts to a click; good is under 200 ms.
Latency	How long something takes to respond. For Walkthru: seconds per step, seconds to the first step, seconds to a finished report.

AI and engineering words


LLM / model	Large language model: the AI that reads the page and decides the next step or writes the report. Walkthru uses free models today (Groq's gpt-oss-120b, Nemotron on OpenRouter, Gemini as backup) and has Claude Haiku 4.5 wired for paid plans.
Tokens	The unit AI providers charge by, roughly three quarters of a word. A Walkthru report uses about 9,000 to 12,000 tokens.
Grounding	Code that throws away any AI-written finding that does not point at a real step where something went wrong. It is how Walkthru stops the AI from inventing problems.
Deterministic	Plain code with a fixed answer, no AI involved. Every SEO, GEO, security, accessibility and email check in Walkthru is deterministic.
MCP	Model Context Protocol: a standard way for AI coding tools like Claude Code and Cursor to call outside services. A Walkthru MCP server lets a developer say "scan my site and fix what you find" without leaving the editor. Planned for Plus.
Webhook	A message one service sends another when something happens, for example Dodo telling Walkthru "this customer paid".
Entitlement	The database record that says which plan a user has and until when. The server reads it on every run; the browser cannot fake it.

Business words

MRR / ARR	Monthly or annual recurring revenue.
ARPU	Average revenue per paying user per month.
CAC	Customer acquisition cost: what it costs, in money or time, to win one paying customer.
Churn	The share of paying customers who stop paying each month.
Conversion	The share of people who move to the next stage, for example from free scan to paid.

Instant Scan

The front door. Anyone pastes a web address on the homepage and gets a full report in seconds, with no account and no install.



The live landing page on the local build. The "Scan my site" button starts an Instant Scan. The hero copy still reflects the older "where strangers get stuck" angle; the new "launch check" positioning lands with task V9.

What happens, step by step

1. The visitor types an address and presses **Scan my site**.
2. Walkthru checks the address is a real public website (not a private network address; this blocks a class of attacks called SSRF).
3. It reads the homepage and follows up to **10 same-site links** within **20 seconds**, respecting robots.txt, using ordinary page requests only.
4. In parallel, plain code runs the SEO, GEO, security-header, accessibility and email checks on those pages.
5. One AI call writes the **first impression** ("what is this, who is it for, what would I click first, how clear is it from 0 to 3"). A second AI call writes the summary and ranks the fixes.
6. The report is saved as a public page with its own link, ready to share.

What it does not do

- It does not click, type or sign up. That needs the extension.
- It does not run the deep security checks (exposed files, leaked keys) because it cannot know the visitor owns the site.
- It does not measure speed unless a Google PageSpeed key is configured (it was not in today's local runs, so speed showed "not measured" rather than a fake pass).

MEASURED TODAY

Hard test site: **26 findings in 11.3 s** (5 high, 9 medium, 12 low), Launch Ready 35, GEO 23.
Showcase test site: **2 findings in 43 s**, Launch Ready 93, GEO 100. The second took longer because the free AI providers were rate-limiting.

GO DEEPER

POST /scans in app/main.py . Limits: 5 scans an hour per IP address (in memory) and
FREE_SCANS_PER_DAY = 200 across all users (counted in the database). Crawler:
app/scans/site.py . Report graph:
app/agent/report.py , which fans out to the scans and joins at synthesize .


The Chrome extension and AI test users

The extension is how Walkthru tests the real product: signup, onboarding, dashboard, settings, checkout up to the payment button. The browser does the clicking on the owner's computer; the thinking happens on Walkthru's server.


How a test run works

1. The owner installs the extension and opens the side panel next to their site.
2. They pick a **goal** (the panel suggests some from the page's own links) and a **test user**.
3. Walkthru's server plans the goal into a checklist of 1 to 4 checkpoints, and refuses unsafe goals ("delete my account", "buy Pro") before a run is even counted.
4. The extension takes a **snapshot** of the page: the address, a numbered list of visible buttons, links and fields, visible text and error messages. Emails, long numbers and key-like strings are masked **before** anything leaves the browser.
5. The server's AI, playing the persona, decides one action: click #12, type into #4, scroll, finished, or give up. It also says what it is thinking and how confused it is (0 to 3).
6. The extension performs the action in the real tab, waits for the page to settle, and captures a screenshot after meaningful actions (at most 8 per run, form fields masked).
7. Repeat until the checklist is done, the test user gives up, a safety rule stops it, or a limit is reached (12 steps free, 30 paid, 4 minutes).
8. The server writes the report in the background and the panel links to it.

A small bird called **Scout** appears in the corner of the tested page while it runs, so the owner can see what the test user is doing. Scout is invisible to the snapshot, so it never confuses the test.



Side panel during a run. This is a design mock with illustrative text, not a recorded run: the real panel only renders inside Chrome.



The four test users

Test user	How the AI is told to behave	What they tend to find	Plans
First-time visitor	Never heard of the product, skims, gets impatient fast	Unclear headlines, hidden signup, confusing first steps	All
Phone user	Small screen, taps big obvious things, hates long forms	Tiny buttons, long forms, broken mobile menus	Paid
Careful buyer	Small-business owner deciding whether to pay; cares about pricing, trust, what happens next	Hidden pricing, missing trust signals, vague plans	Paid
Skeptic	Cautious developer who reads error messages, checks links, distrusts vague copy	Silent errors, dead links, marketing fluff	Paid

Plus will add **custom test users** ("a 60-year-old first-time buyer on a phone") and several test users merged into one report.

GO DEEPER


Extension: Chrome Manifest V3 built with WXT and React (`apps/extension`). `lib/snapshot.ts` numbers elements (open shadow roots included) and marks fields as filled, empty or checked, never their values. `lib/redact.ts` masks emails, long numbers and tokens. `lib/execute.ts` runs one action with the safety filter. The step loop lives in the side panel page, not the service worker, because Chrome suspends workers after about 30 s idle. Server: a LangGraph graph (`app/agent/persona.py`) with `decide` → `interrupt` → `check` ; each step is one HTTP call to `POST /runs/{id}/observe` . Goal planner: `app/agent/goal.py` . Loop guards: three identical steps end as `stuck` , a third visit to the same page ends as `looping` (Walkthru's own stop, never a site finding). "Done" right after an action that changed nothing is questioned, then ends as `gave_up` . Permissions: `activeTab`, `sidePanel`, `scripting`, `tabs`, and an optional all-sites permission requested once at Start, used only for screenshots of the test tab.

Why the browser runs on the owner's machine

- **Logged-in pages just work.** The owner is already signed in, including magic-link, Google sign-in and two-factor accounts. No password or cookie is handed to Walkthru.
- **Hosting stays almost free.** No fleet of cloud browsers to pay for.
- **The cost:** an install step. Every competitor starts from a URL. This is the biggest funnel risk, and it is measured by experiment A2 (chapter 17). If fewer than 25% of beta users finish a run, we build a **cloud runner** for public pages.

The report, section by section

Everything Walkthru does ends in one report page. The screenshots below are a real Instant Scan of our "hard" test site (Zentrix), a fake startup page seeded with 22 known problems, generated today on the local build.




The site being tested: the Zentrix fixture. Note the vague headline, two identical buttons and a signup link hidden in the footer.

1. Header and summary

The tested address, the date, the test user, and **which AI models ran** (every report states this). Then a plain-language paragraph summarising what was found, and three quick numbers: total findings by severity, first-impression clarity (0 to 3), and how deep the security scan went.

2. Launch Ready score and badge

The single number, with a bar for each area. Areas that were not measured are shown as "Not measured" and their weight moves to the others, so the score never punishes or rewards what was not checked. Below it, the badge with copy-ready HTML and Markdown.



Summary, Launch Ready score (35, Needs work) and the badge. An Instant Scan has no journey, so "Can a stranger get in" is Not measured.

3. First impression, five seconds in

What the site does, who it is for, what a visitor would click first, and which trust signals they noticed. This is the classic "5-second test" that UX researchers run with real people.

4. AI search readiness (GEO)

The GEO score out of 100 with its seven categories, then **"What AI search sees"**: the literal text an AI crawler receives before JavaScript runs, with a word count. Seeing "37 words" of buzzwords is usually the moment a founder understands the problem. Below it, the **GEO fix pack** with copy buttons (free shows one fix and says how many more exist).

First impression, five seconds in

What this site does

A platform called Zentrix that provides modern synergy and next-generation workflow tools.

Who it is for

Businesses and teams that want to streamline collaboration at scale.

What they would click first

Log in, it's the most prominent action to access the platform.

Trust signals

© Zentrix Inc., Legal link

AI SEARCH READINESS (GEO)

Can AI search read and quote you?

23

of 100

Critical

Readiness, not rankings: whether assistants can fetch, understand and quote your pages. Fixes are in the findings below, marked GEO.

AI crawler access

10/25

Content before JavaScript

8/15

Structured data

0/15

Easy to quote

0/15

Name and trust

1/10

Title and description

3/10

llms.txt

0/5

What AI search sees on your homepage, before JavaScript runs: 37 words

Zentrix Pricing Log in Zentrix The platform for modern synergy Leverage next-generation workflows to unlock outcomes at scale. Continue Proceed Subscribe to our newsletter for synergy tips. Subscribe Close © Zentrix Inc. · Legal · Create account

GEO fix pack

Let AI search assistants read your site

robots.txt

Add at the top of robots.txt at your site root. To opt out of AI training only, block GPTBot, Google-Extended and CCBot instead.

```
# Let AI search assistants read and cite your site
User-agent: OAI-SearchBot
Allow: /

User-agent: ChatGPT-User
Allow: /

User-agent: Claude-SearchBot
Allow: /

User-agent: ClaudeBot
Allow: /
```

Copy

3 more ready-made fixes for this site on Pro.

robots.txt blocks AI training crawlers (GPTBot, Google-Extended, CCBot, anthropic-ai). That is a choice and does not affect AI search answers.

Not blocked for an AI search user agent. Edge networks verify real crawlers by IP, so this cannot prove access.

CRAWL COVERAGE

1 page checked

Page limit

10

Duration

32 ms

robots.txt

Respected

Scope

Same origin

Only this page was found. The HTML the server sends links to no other pages on this site, which is typical when navigation is built by JavaScript.

01 http://127.0.0.1:8182/

First impression, GEO score 23 (Critical) with category bars, the exact text AI search sees, the first fix from the fix pack (robots.txt rules) and crawl coverage.

5. Crawl coverage

Which pages were checked, the page limit, how long it took, and whether robots.txt was respected. Honest about what was **not** checked, so a clean report never falsely implies the whole site is fine.

6. Technical launch checks and "Fix these first"

TECHNICAL LAUNCH CHECKS

Catch what the journey cannot see

Passive public checks only

2 issues

Accessibility basics
Page language, heading order, image text and named controls.
1 high priority

not measured

Mobile performance
PageSpeed score and Largest Contentful Paint evidence.
Unavailable for this run

8 issues

Search readiness
Titles, descriptions, headings and crawl basics across audited pages.
2 high priority

7 issues

Security hygiene
Passive page checks plus verified-domain exposure checks.
1 high priority

PRIORITIZED FIXES

Fix these first

P01

Serve the site over HTTPS to secure data transmission.

P02

Add a meta description and lengthen the page title to improve SEO.

P03

Remove the robots.txt block or allow search engine bots to crawl the site.

P04

Add a Content-Security-Policy header and set HttpOnly and SameSite on session cookies.

P05

Provide alt text for all images and give the form control an accessible name.

Four check cards (accessibility, performance, search, security) and the five prioritised fixes.

7. All findings

Every finding with its severity, category, a one-line explanation of why it matters, the fix, and the evidence (the URL, the header, the image file names). Repeated issues across pages are merged into one finding that names how many pages and which ones.

All findings	
high Accessibility	Form control has no accessible name A screen reader cannot tell users what this control does. Fix: Add a visible label or an accurate aria-label to every control. <input>
high SEO	Missing meta description Search engines will pick a random snippet of page text as the summary. Fix: Add a meta description of 50 to 160 characters that says what the site does and for whom. http://127.0.0.1:8182/
high SEO	robots.txt blocks all crawlers User-agent: * with Disallow: / hides the whole site from search engines. Fix: Remove the blanket Disallow or limit it to private paths. Disallow: /
high SEO	robots.txt blocks AI search assistants Your robots.txt tells these assistants not to read your site: OAI-SearchBot, ChatGPT-User, Claude-SearchBot, ClaudeBot, PerplexityBot, Googlebot, Bingbot, Applebot. They cannot quote or link to pages they may not fetch. Fix: Allow them in robots.txt (User-agent: <name> / Allow: /). To opt out of AI training only, block GPTBot, Google-Extended and CCBot instead. http://127.0.0.1:8182/robots.txt
high Security	Site is served over plain http Everything visitors type, including passwords, crosses the network unencrypted. Fix: Get a certificate (Let's Encrypt is free) and redirect http to https. http://127.0.0.1:8182/
medium Accessibility	2 images are missing text alternatives Meaningful images without alt text are not announced to screen-reader users. Fix: Add descriptive alt text, or alt="" when an image is decorative. hero1.svg, hero2.svg
medium SEO	Page title is too short The title is "Home" (4 characters). Fix: Use 15 to 60 characters that describe the page and the brand. Home
medium SEO	2 of 2 images have no alt text Images without alt text are invisible to search engines and screen readers. Fix: Describe each meaningful image in its alt attribute; use alt="" for decorative ones. hero1.svg, hero2.svg
medium UX	Sign-up is hidden in the footer The only sign-up link on the homepage is "Create account" in the footer. Visitors look for it in the header or the first screen, and many leave before scrolling that far. Fix: Add a clear Sign up button to the header and to the first screen of the homepage. http://127.0.0.1:8182/: footer link "Create account"
medium SEO	The homepage says little before JavaScript runs AI crawlers read 37 words on your homepage. That is thin for an assistant to explain what you do and for whom.

The findings list. Each item names the evidence underneath, for example `Disallow: /` or `hero1.svg, hero2.svg`.

8. Only in test-run reports (extension)

- **Journey replay:** a three-pane timeline of every step with the screenshot, what the test user thought, confusion level, and per-step accessibility and speed numbers. Play, pause and step controls.
- **Outcome:** reached the goal, gave up, stuck, looping, safe stop, or ran out of steps.
- **Since your last run** (paid): fixed, still broken, new, not re-checked, page by page.
- **Agent fix prompt** (paid): Copy, Download `walkthru-fixes.md`, or Copy chat version for Lovable and Bolt.
- **Ignore a finding** (paid): mark "won't fix" with a reason. It leaves the lists but still counts in the score, so the badge cannot be gamed.

The services, one by one

What Walkthru actually delivers, in the order a founder meets it.

7.1 Journey testing (UX)

AI test users try the real flow and report where they hesitated or failed, with the step, the screenshot and their own words. A UX finding is only allowed into the report if it points at a step where something actually went wrong: confusion of 2 or more, a visible error, a click that changed nothing, giving up, or a loop. A journey where everything worked has **no** UX findings, instead of padding.

7.2 SEO

Checked on up to 10 pages (Instant Scan and free) or 50 pages (paid): page title length, meta description, one H1 heading, canonical link, mobile viewport, language, accidental "noindex", image alt text, Open Graph tags for link previews, robots.txt and sitemap. Pages the test user visited are audited too.

7.3 GEO (AI search readiness)

Scored 0 to 100 in seven categories. All deterministic, so it costs nothing in AI calls and is identical on every plan.

Category	Points	What is checked
AI crawler access	25	robots.txt allows the AI search bots (OAI-SearchBot, ChatGPT-User, Claude-SearchBot, ClaudeBot, PerplexityBot, Googlebot, Bingbot, Applebot); the homepage answers an AI user agent normally instead of blocking it
Content before JavaScript	20	The homepage HTML contains real text before JavaScript runs; key pages are not empty shells
Structured data	15	Valid JSON-LD, Organization or WebSite, a page-type schema (SoftwareApplication, Product, FAQPage, Article), enough useful properties
Easy to quote	15	One H1, descriptive H2s, a plain answer paragraph near the top, lists or tables, concrete numbers
Name and trust	10	Same name in title, og:site_name, schema and H1; about, contact, privacy or pricing links; public profile links
Title and description	10	Reuses SEO signals (findings stay in SEO, so nothing is reported twice)
llms.txt	5	Present and well formed; labelled "low measured impact"

Bands: 0 to 35 critical, 36 to 67 foundation, 68 to 85 good, 86 to 100 excellent. Blocking *training* bots is reported as a choice, not a fault. The user-agent probe is worded carefully: a 200 response means "not blocked by user agent", because real edge networks verify bots by IP and we cannot prove access.

Fix pack: robots.txt rules, Organization, WebSite and SoftwareApplication JSON-LD filled from the site's own title and description, an llms.txt draft listing the audited pages, and a rendering fix picked by detecting the framework (Vite single-page app, Next.js, Lovable, Astro).

7.4 Security hygiene (passive only)

Everyone, any site	Only on a domain the owner verified
HTTPS and redirect to HTTPS; security headers (CSP, X-Frame-Options, X-Content-Type-Options, Referrer-Policy, HSTS); cookie flags; server version leaks	Exposed files (<code>.env</code> , <code>.git</code> , backups); secret keys inside public JavaScript bundles, reported with the value masked

Never: login brute force, injection attempts, port scans, load tests. The reason is legal and ethical: probing a site you do not own is an attack.

7.5 Accessibility and speed

Inside a journey, the extension runs **axe** WCAG A/AA checks on the real rendered page and records LCP, CLS and INP at each step, so a finding can say "at step 3 on /signup, the email field has no label". Server-side static checks and Google PageSpeed are fallbacks. When something cannot be measured, the report says "not measured", never a false pass.

7.6 Signup email check

Reads the domain's SPF and DMARC records (free) and MX records (paid) through DNS over HTTPS. It also recognises a common trap: Supabase's built-in mailer allows very few emails per hour, so a journey error "email rate limit exceeded" becomes the fix "set up your own SMTP".

7.7 Agent fix prompt ("prompts")

The paid feature that closes the loop. Walkthru turns the report into one instruction for the founder's AI coding tool. It is built by plain code from the stored report, **never by an AI**, so it cannot invent findings, and it costs nothing to make.

- **Contains:** the detected tech stack, then every finding in priority order (security, UX blockers, GEO, SEO, accessibility, performance), each with what is wrong, where, the evidence, the required change and an acceptance check. Every affected page is listed.
- **Rules for the coding agent:** keep existing behaviour, ask before adding dependencies, never paste secrets.
- **Last step:** "rerun Walkthru to verify", with the list of findings expected to flip to fixed.
- **Formats:** full Markdown for Cursor, Claude Code and Codex; a short chat version for Lovable and Bolt; a `walkthru-fixes.md` download.
- **Tested end to end by the founder:** a real prompt applied to the showcase site, rerun compared correctly (2 fixed, 0 new, 3 still broken).

7.8 Badges

The owner can embed a live "Walkthru | Launch Ready 92" badge in their footer or README. It links to the latest public report for that site and refreshes (cached one hour) when they rerun and share again. Ignored findings still count, so the badge cannot be gamed. For the owner it is proof of quality; for Walkthru it is a free, permanent backlink and ad on every site that uses it. This is the main built-in growth loop.

GO DEEPER

`app/agent/score.py` computes the score in code from the stored report (weights UX 30, security 20, GEO 20, SEO 15, speed and accessibility 15; penalties per finding high 25, medium 10, low 3). `GET /badge/{run_id}.svg` and `GET /badge/{run_id}` in `app/main.py`. Status: built in the working copy on 24 September and visible in today's reports; not yet committed, so it counts as "in progress" (task V8).

7.9 How founders validate their own product

1. **Baseline:** run an Instant Scan and one journey before launch.
2. **Fix:** paste the fix prompt into the coding tool, or apply the fix pack by hand.
3. **Prove it:** rerun the same goal. The comparison shows, page by page, what is fixed, what is still broken, what is new and what was not re-checked. A finding only counts as fixed if its pages were actually checked again.
4. **Show it:** share the public report and put the badge on the site.
5. **Keep it (Plus, planned):** weekly watch and a deploy hook email them only when something breaks or gets fixed.

Verifying the domain (meta tag or file) is also how they prove ownership, which unlocks the deeper security checks and one real form send.

Reading the numbers

Every number in a report, and every performance number in this document, explained once.

Number	What it describes	How to read it
Launch Ready score (0 to 100)	Weighted mix of measured areas: UX 30, security 20, GEO 20, SEO 15, speed and accessibility 15	85+ Launch ready, 60 to 84 Almost ready, under 60 Needs work
Area bars (0 to 100)	Each area starts at 100 and loses 25 per high, 10 per medium, 3 per low finding. GEO uses its own score. UX starts at 100 if the goal was reached, 50 if the test user gave up or got stuck	Low bar = where to spend time first
GEO score (0 to 100)	Points earned in the seven GEO categories	0 to 35 critical, 36 to 67 foundation, 68 to 85 good, 86+ excellent
Clarity (0 to 3)	How quickly the AI understood what the site is from the first screen	3 = instantly clear, 0 = could not tell
Confusion (0 to 3, per step)	How lost the test user felt at that step	2 or more makes the step eligible for a UX finding
Severity	High: blocks users, leaks data or hides the site. Medium: hurts conversion or ranking. Low: polish	Fix high first; the fix prompt orders them for you
"AI search sees N words"	Words in the homepage before JavaScript runs	Under about 100 words usually means a thin page or a JavaScript shell
LCP / CLS / INP	Main content load time / layout jumping / click response	Good: under 2.5 s / under 0.1 / under 200 ms
Pages checked	Pages crawled vs the plan limit (10 or 50)	Anything not listed was not checked
Tokens	AI usage for the run	About 9,000 to 12,000 per report; a journey of 4 to 6 steps used 6,400 to 11,800

Speed of Walkthru itself (latency)

Measure	Today	Notes
Instant Scan, small site	11 to 43 s	Target is 20 s or less for 10 pages. Slower when free AI providers rate-limit
First step of a journey	7.6 to 8.8 s	Was 13 to 19 s; cut by running domain verification in parallel and warming models at startup
Server time per later step	1.4 to 6 s	Steps arrive about 5 s apart including the browser acting and the screenshot upload
Whole short journey	25 to 33 s	4-step signup on the easy site; portfolio contact goal in 33 s
Report after a journey	+ up to ~1 min	The careful free report writer (Nemotron Ultra) is slow but writes the most accurate summaries
50-page paid crawl	13.9 s	python.org, 50 pages, measured

Quality numbers

- **Trap recall 82%:** our hard test site hides 22 known problems (6 UX, 6 SEO, 6 security, 4 GEO). Two real journeys found 18. By area: SEO 6 of 6, security 6 of 6, GEO 4 of 4, **UX only 2 of 6**. The machine checks are strong; judging UX from a journey is the weakest part and the main place Claude on paid plans should help.
- **Invented findings:** in a test of seven free models writing reports, the grounding code removed every invented UX finding from every model.

Safety: how the agent contains itself

An AI clicking around inside someone's logged-in account is dangerous if it is careless. Walkthru's limits are written in code, in two places (server and extension), not left to the AI's good behaviour.

Risk	Rule, enforced in code
Paying, deleting, cancelling	Buttons with words like pay, buy, purchase, checkout, delete, remove, cancel subscription, transfer, unsubscribe are never clicked, on any page . On logged-in pages the agent does not even type into them. The run ends as a safe stop at that button.
Sending a form or message to a real person	Send and invite buttons run only on a domain the signed-in owner has verified, only after the owner clicks confirm in the side panel, and at most once per run . Everywhere else the run ends as "Stopped before sending". (This rule exists because a portfolio contact form was submitted once during early testing.)
Anonymous or public forms on sites the user does not own	Same as above: without verification, submit is never pressed. The agent fills the fields to prove the form works, then stops.
Leaving the site	Same-origin only . If a click leads to another domain (a new scheme, host or port), the extension stops the run and closes it with a reason. External links and new tabs are reported, not followed. <code>mailto:</code> and <code>tel:</code> links are reported as working contact methods and never opened.
Unsafe goals	The goal planner refuses goals that ask to pay, delete, cancel, spam, attack or visit another website, with an error, before a run is counted. Checked live: "delete my account and buy pro" was refused in 0.5 s.
CAPTCHAs and bot checks	The snapshot flags a CAPTCHA and the agent stops. It never tries to solve one.
Endless loops	Three identical steps end as stuck; a third visit to the same page ends as looping; 12 steps (free) or 30 (paid); 4 minutes total.
Personal data	Emails, long numbers and key-like strings are masked in the browser before upload. Field values are never sent, only "filled" or "empty". Screenshots mask form fields and hide Scout. Signups use a fake test identity.
Attacking sites	Security checks are passive only. Deep checks need domain verification. The scanner blocks private network addresses (SSRF protection) at every redirect.
Blaming the site for Walkthru's own stops	Safe stops, loops, the owner pressing Stop and CAPTCHA stops are Walkthru's limits and are never reported as site problems.
Data kept forever	Screenshots expire after 30 days. Owners can delete any run, export all their data, or delete their account.

Honest limits of the agent

- It cannot test payment itself, only the path up to the pay button.
- It cannot pass CAPTCHAs, SMS codes or email verification links.
- With 12 free steps it covers short flows only; long onboarding needs a paid plan.
- On free models, UX judgement is uneven (see chapter 11). It finds concrete breakage well, and subtle design confusion less reliably.
- It tests one persona per run today; merged multi-persona reports are planned for Plus.
- It needs Chrome on a desktop. No Firefox, Safari or native mobile apps.

GO DEEPER

Word lists in `app/agent/safety.py` (`DESTRUCTIVE` , `SENDING`) are mirrored in `apps/extension/lib/execute.ts` ; the server's `persona._enforce` and the extension's `executor` both refuse, so a buggy or manipulated model cannot bypass one layer. Same-origin compares scheme and host. Domain verification: `GET /verification` returns a per-account token for a meta tag or `/.well-known/walkthru.txt` .

The plans, in full

Prices confirmed by the founder on 24 September 2026. The founding price applies to the first 50 customers, locked for 12 months. At launch, paid access is sold as founder-approved 30-day passes through Dodo, with no automatic renewal (chapter 15 explains why).

	Free	Launch Pack	Pro	Plus
Price	\$0	\$9 once	\$19/mo (\$15 founding)	\$49/mo (\$39 founding)
AI for test user and report	Free models	Claude Haiku 4.5	Claude Haiku 4.5	Claude Haiku 4.5
Instant Scan	5 an hour, 10 pages	Yes	Yes	Yes
Test runs	3 a month	20 in 30 days	40 a month	150 a month
Steps per run	12	30	30	30
Pages the test user may enter	Public only	Public + logged-in	Public + logged-in	Public + logged-in
Test users	First-time visitor	All 4	All 4	All 4 + custom
Sites per month or pass	1	1	2	5
GEO score	Homepage, top 3 fixes	Full site	Full site, 50 pages	Full site, 50 pages, every site
"What AI search sees"	Yes	Yes	Yes	Yes
GEO fix pack	1 fix preview	Yes	Yes	Yes
Agent fix prompt	Locked (shows count)	Yes	Yes	Yes
Launch Ready score and badge	Yes	Yes	Yes	Yes
SEO	10 pages	50 pages	50 pages	50 pages
Security	Headers, TLS, cookies	+ exposed files, leaked keys (verified)	Same	Same
Signup email check	SPF, DMARC	Full	Full	Full
Rerun and compare	No	Yes	Yes	Yes
Ignore a finding	No	Yes	Yes	Yes
One real form send (verified, confirmed)	No	Yes	Yes	Yes
Signup funnel numbers planned	No	Yes	Yes	Yes
Landing copy review planned	No	Yes	Yes	Yes
Competitor side by side planned	No	3 competitors	3 competitors	3 per site
Walkthru MCP server planned	No	No	No	Yes

Weekly GEO/SEO watch + deploy hook planned	No	No	No	Yes
Report export	Web report	PDF	PDF	PDF with your logo
Shareable public report	Yes	Yes	Yes	Yes

Who each plan is for

- **Free:** the hook. "Can ChatGPT read your site?" in 20 seconds, plus three public-page journeys a month. Upgrade trigger: "the test stopped at my login screen" or "I fixed it, did it work?"
- **Launch Pack (\$9 once):** the founder launching next Tuesday. Everything paid for 30 days, 20 runs. It is also the decoy: someone who expects to ship again sees that \$19 a month for 40 runs and ongoing comparisons is better value.
- **Pro (\$19):** the solo founder who keeps shipping. Logged-in flows, reruns, the fix prompt, 2 sites.
- **Plus (\$49):** the agency or multi-product founder. 5 sites, 150 runs, white-label PDFs, weekly watch, MCP. Launches as a **waitlist** until watch and branded PDFs exist, so nothing is sold before it exists.

MISMATCH TO FIX BEFORE LAUNCH

The landing page still shows the old plans (Pro \$15, Team \$39, 60 and 250 runs). The new table above lands with task V9. The docs page also still describes Instant Scan without GEO.

Free to start. Pay when you test the real product.

Free	Launch Pack	Pro Most picked	Team
\$0 forever Instant Scan, no install 3 test runs a month Public pages only Scan my site	\$9 one time 20 runs within 30 days Everything in Pro Buy the pack	\$15 per month 60 runs a month Logged-in pages Full-site SEO Full security check Compare runs Start Pro	\$39 per month 250 runs a month 5 sites Custom test users Weekly scans Start Team

Current pricing section on the landing page: old plans, to be replaced.

GO DEEPER

The server owns every limit: `app/plans.py` holds the plans; `POST /runs` ignores whatever tier the browser claims and returns 403 (logged-in or persona not allowed), 402 (runs used up) or 429 (daily free capacity used). Checked live: a free account claiming to be paid was refused. `GET /me/plan` feeds the side panel. `scripts/grant_plan.py` EMAIL `pro` grants a free test pass. Claude Haiku 4.5 is wired through Google Cloud and switches on when `CLAUDE_VERTEX_PROJECT` is set; until then paid runs also use the free chain.

Is the free report real? An honest audit

Short answer: yes, it is a real report, not a placeholder. About 80% of it is solid, production-grade checking. The AI-written 20% is useful but uneven on free models, and that is exactly the part paid plans improve.

What is genuinely strong

- **Every machine check is deterministic and correct by construction.** On today's hard-site scan all 26 findings were real and each named its evidence: `Disallow: /` in robots.txt, a title of 4 characters, the two image file names, the cookie called "session" missing HttpOnly and SameSite. These are the same facts Lighthouse or securityheaders.com would report, in plain language, in one place, with the fix.
- **GEO is the standout.** "What AI search sees" with a word count is concrete and rare. The showcase site, built correctly, scored 100; the broken one scored 23. The score moves for the right reasons.
- **Grounding works.** A journey where nothing went wrong produces no invented UX problems. Checkpoint B found 82% of seeded traps.
- **It is honest about gaps:** "Not measured", "before JavaScript runs", pages not checked, which models ran.

What is weak today (seen in today's scans)

- **AI judgement on free models is uneven.** The showcase homepage says clearly "Meeting notes that write themselves ... teams of 3 to 50 people". The first impression correctly described what it does and who it is for, yet rated clarity 0 of 3 and raised a **high** "Value proposition unclear" finding. That is a wrong call, and a founder would notice.
- **Local-site noise:** the same scan flagged plain http as high on a local test address while its own summary said that is expected locally. Real sites are unaffected, but it shows the writer can contradict itself.
- **Speed was "not measured"** in both scans because no PageSpeed API key is configured locally. Set the key before launch or the speed area stays empty.
- **Journey UX recall is 2 of 6** on the hard site. Free journeys are 12 steps on public pages with one persona, so they find blocking breakage (silent failed signup, dead button, hidden signup link) far better than subtle confusion.
- **Fixtures are not customers.** All quality numbers come from our own test sites and the founder's projects. No outside founder has used it yet.

Is it a big deal for a founder?

Part of the free report	Worth to a real founder	Could they get it free elsewhere?
GEO score + what AI sees + 1 fix	High. Most founders have never checked it, and SPAs usually fail	Partly, via separate GEO tools, mostly paid
Security headers, cookies, HTTPS	Medium. Quick wins, good hygiene	Yes (securityheaders.com, Mozilla Observatory)
10-page SEO audit	Medium. Catches embarrassing mistakes like a blanket robots block	Yes (Lighthouse, many SEO checkers)
First impression	Medium, when correct; see the clarity miss above	Not in one click

3 public journeys a month	High when it catches a broken signup; low for a flow that works	Competitors give more free runs (Meerkat about 170 in month one)
All of it in one link with a score and badge	The real value. One place, plain words, shareable	No single tool combines these

VERDICT

The free report is a credible, shareable product on its own, strongest on GEO and technical hygiene. It will not wow a founder whose site is already healthy, and a wrong AI judgement on clarity could cost trust. **Before launch:** tune the first-impression clarity rule so a finding requires the clarity score and the description to agree, suppress http findings on local addresses in scans, configure PageSpeed, and run the measured Claude comparison (about \$1 to \$2) before claiming paid reports are better.

Who benefits, and how

Four people, four stories. Details are illustrative, not real customers.

The indie hacker, launching on Product Hunt next week

Built a habit tracker with Lovable over three weekends. Runs a free Instant Scan: GEO 18, "AI search sees 6 words". The app is a JavaScript shell. Buys the \$9 Launch Pack, runs the careful-buyer and phone-user journeys through signup and the first dashboard, finds the signup confirmation email never arrives (Supabase mailer limit), pastes the fix prompt into Lovable, reruns: 7 fixed, 1 new. Puts the Launch Ready badge in the footer before launch day.

The small business founder

Runs a booking site built by a freelancer. Not technical. Uses Instant Scan monthly; the plain-language report tells them what to ask the freelancer, and the fix prompt is something they can forward. The security section catches the login cookie without HttpOnly.

The startup founder with a live product

Two founders, 300 users, shipping every week. Pro plan. After each release they rerun "sign up and create a project" as a skeptic on staging, logged in, without handing any password to a vendor. The comparison tells them within minutes whether the release broke onboarding. Later, the MCP server lets their Claude Code agent run the check itself before each deploy.

The small agency


Hands 3 to 5 sites a month to clients. Plus plan. Every handover includes a white-label PDF with their own logo showing the Launch Ready score, SEO, GEO and security. Weekly watch emails them if a client's developer later blocks AI crawlers by accident. It is a reason to sell a maintenance retainer.

The outcomes they buy

- Fewer lost signups on launch day.
- Being findable in Google and in AI answers.
- No embarrassing leaked key or open file.
- Proof (report, badge, comparison) they can show users, clients or investors.
- Minutes instead of hours: the fix prompt does the typing.

How it is built

Three parts: a website, a Chrome extension, and a server that holds the brain. A hosted database keeps accounts and reports. The design goal is near-zero running cost until there is revenue.



Part	Technology	State
Web	React 19, Vite, TypeScript, Tailwind v4, GSAP animation	Built: landing, login, dashboard, reports, docs, privacy, terms, security pages
Extension	Chrome MV3, WXT, React, axe-core	Built; not yet on the Chrome Web Store
API	Python, FastAPI, LangGraph	Built; runs locally, not deployed
Database, auth, files	Supabase (Postgres with row-level security, magic link, Google and GitHub sign-in)	Live on the free tier; OAuth providers still to enable
AI models	Free chain: Groq gpt-oss-120b, gpt-oss-20b, Qwen, then Gemini; Nemotron 3 Ultra on OpenRouter writes most reports. Paid: Claude Haiku 4.5 on Google Cloud	Free chain live; Claude wired, switched off
Payments	Dodo Payments	Designed (payment.md), not built
Hosting	Vercel (web), Oracle Always Free VM or a ~\$5 VPS (API)	Not started
Evals	Fixture sites with seeded traps, a trap scorer, LangSmith tracing	Built

Principles that keep it cheap and honest

- **Plain code wherever possible.** AI only plans goals, plays the test user and writes prose. Every check is deterministic.
- **Browser on the user's machine, brain on our server.** No cloud browsers to pay for.
- **Text snapshots, not screenshots, for decisions.** Tokens are the main cost.
- **The server decides the plan.** Nothing the browser says about its tier is trusted.
- **Grounding and proof before "done".** The AI cannot claim success after an action that changed nothing.

GO DEEPER

Repository layout: `apps/web` , `apps/api` , `apps/extension` , `evals` . The API reaches Supabase over its HTTPS Data API because direct Postgres was unreliable from the founder's network. Agent state is in memory by default;

`CHECKPOINTINTER=postgres` for production. Free-provider ceiling: Groq allows about 8,000 tokens a minute per model and about 1,000 requests a day, so free journeys are capped at `FREE_RUNS_PER_DAY = 90` and scans at 200 a day.

Tests: 171 API, 30 extension; builds and lint green.

Market and competitors

"AI users test your site" is a crowded, young category. None of the players has won. Walkthru's opening is the combined launch check and logged-in testing without credentials.

Market size, roughly

- Vibe-coding tools market: about **\$4.7B in 2026**, growing about 38% a year (industry estimates, sources in the appendix).
- Lovable: about 8 million users (late 2025) and 100,000+ new projects a day; \$400M annual revenue (Feb 2026). Cursor: \$2B+ annual revenue (Feb 2026). Replit: \$9B valuation (Mar 2026).
- **Estimate of reachable buyers:** if 1% of Lovable's 8 million users ship a public app a year, that is 80,000 launches a year from one builder alone. At a 1% paid conversion, 800 customers, about \$20,000 MRR at \$25. Real reach will be smaller; the point is that the pool is not the constraint, distribution is.

Direct and indirect competitors (prices checked 24 Sep 2026)

Name	Price	What they sell	Threat
CanaryUsers	Free 100 credits; \$9, \$19, \$49/mo	A "flock" of AI users on every deploy; logged-in flows from \$9	high same buyer, same price
Meerkat	Large free tier; \$49, \$199/mo	AI personas in a real browser; findings across releases	high generous free tier
Swarm	5 free runs; \$150/mo	Feedback in 10 minutes, code fixes; VC-backed	medium different buyer
Uxia	1 free test, custom	AI testers for prototypes and live sites	low
Octomind (indirect)	\$89 to \$589/mo	Automated QA test generation	low
Otterly.AI (indirect, GEO)	From \$29/mo	AI visibility tracking and GEO audits	low never tests flows

Feature	Meerkat	CanaryUsers	Swarm	Otterly	Walkthru
AI users in a real browser	Yes	Yes	Yes	No	Yes
Logged-in flows	Yes	Yes	Yes	No	Yes (paid)
No credentials shared	Unknown	Unknown	Unknown	n/a	Yes
SEO	Unknown	Unknown	Unknown	Partial	Yes
GEO	Unknown	Unknown	Unknown	Yes	Yes
Passive security	Unknown	Unknown	Unknown	No	Yes
Compare across releases	Yes	Partial	Unknown	No	Yes (paid)
Runs on deploy	Yes	Yes	Unknown	No	Planned (Plus)
Several personas per run	Yes	Yes	Yes (10)	No	1 today; planned
Free tier	~170 runs month one	~2 deep scans	5 lifetime	n/a	Instant Scan + 3 runs

"Unknown" means the competitor's public pages do not say. Meerkat's own comparison article claims no persona tool lists SEO or security; treat it as their claim.

Where Walkthru can win

1. **One launch-readiness report** (journeys + SEO + GEO + security). Head start estimate: 2 to 4 months; a competitor could bolt on checks quickly, so the lead comes from owning the message first.
2. **Logged-in testing without a password.** Head start estimate: 6+ months (needs an extension, store review and a safety model). Matters most for agencies and apps with SSO or 2FA.
3. **Reports that never blame the site for the tool's own failures.** A quality bar more than a headline.

Idea scorecard (founder skills review, 24 Sep)

Dimension	Score /5	Why
Problem severity	3	Real, but peaks around launch
Market size	4	Millions building with AI, plus agencies
Willingness to pay	2	Price-sensitive buyers; competitors anchor at \$9
Competition gap	2	Four direct competitors; gap is the combined report
Distribution	3	Instant Scan is a strong hook; extension is friction
Timing	4	AI-built apps and AI search both peak now
Founder fit	3	Proven builder; little sales experience. The co-founder closes this gap
Total	21/35	Promising with gaps. Prove people will pay before building more

WHAT THIS MEANS

The technology is not the risk. **Distribution and a sharp reason to pay are.** That is precisely the job this document asks a co-founder to own. Realistic ceiling as an indie product: about 2,000 customers × \$25 × 12 = **\$600K a year** (Estimate). Going beyond that most likely runs through agencies and white-label reports.

Business model and unit economics

Freemium: a genuinely useful free tier as the demo and the marketing, flat paid plans with a run allowance, and a one-time pack for people who launch once.

How money flows

1. A customer pays Dodo Payments (a merchant of record, so Dodo handles card processing and sales tax).
2. Dodo keeps **4% + \$0.40** per transaction and pays out twice a month (minimum about \$50).
3. Walkthru pays its own AI provider (Google Cloud for Claude) separately, prepaid by UPI.

Offer	Price	Dodo fee	Net	Runs	AI budget cap (25% of net)	Worst case at \$0.20/run
Launch Pack	\$9.00	\$0.76	\$8.24	20	\$2.06	\$4.00
Pro, founding	\$15.00	\$1.00	\$14.00	40	\$3.50	\$8.00
Pro	\$19.00	\$1.16	\$17.84	40	\$4.46	\$8.00
Plus, founding	\$39.00	\$1.96	\$37.04	150	\$9.26	\$30.00
Plus	\$49.00	\$2.36	\$46.64	150	\$11.66	\$30.00

Measured AI cost with Claude Haiku 4.5 (from real run sizes at about \$1.10 per million input and \$5.50 per million output tokens on Google Cloud, to be confirmed at setup): a typical run costs **\$0.02 to \$0.04**, a full 30-step run about \$0.08. Pro used fully costs about **\$1.40 a month** (worst \$3.08) against \$17.84 net. Plus about \$5.43 (worst \$11.73) against \$46.64. Both inside the 25% cap. GEO and the fix prompt cost nothing, being plain code.

Margins (Estimate)

Plan	Net	Likely AI cost	Gross margin
Launch Pack	\$8.24	~\$0.60	~85%
Pro	\$17.84	~\$1.40	~87%
Plus	\$46.64	~\$5.43	~84%

Fixed cost is about **\$6 a month** (a \$5 VPS if Oracle's free VM does not work, plus the domain). **One Pro customer covers hosting.** A \$1,000 a month founder income needs about **64 Pro** or **26 Plus** customers. At 80% Pro and 20% Plus the blended ARPU is **\$25** at list price, about **\$20** at founding prices.

Why paid access starts as founder-approved passes

The founder has no credit card and cannot risk an unknown AI bill before revenue. So at launch: a user requests paid access, the founder approves, the user gets a private one-time checkout link, and a verified Dodo webhook activates a **30-day pass** with a fixed number of runs and a hard AI cost cap. No auto-renewal, no overage, no surprise bills in either direction. Self-serve monthly subscriptions open only after 10 reconciled passes, 30 days of real cost data and a tested spend cap.

Revenue scenarios (all Estimates, no customers yet)

Assumes the first 50 customers pay founding prices (blended \$19.80) and later customers \$25; excludes one-time Launch Packs; counts paying customers at month end after churn.

Scenario	Month 3 (Jan 2027)	Month 6 (Apr 2027)	Month 12 (Oct 2027)
Conservative	10 paying, ~\$200 MRR	25, ~\$500	60, ~\$1,240
Base	25, ~\$500	70, ~\$1,490	200, ~\$4,740
Optimistic	50, ~\$990	150, ~\$3,490	450, ~\$10,990

What the base case requires: at a 3% scan-to-paid conversion, 200 customers need roughly 6,700 Instant Scans over the year, about 560 a month. The free-capacity cap (200 scans a day, about 6,000 a month) supports that. Validation gate before anything else: **5 paid out of the first 100 scans.**

Pricing rules already decided

- Founding prices for the first 50 approved customers, locked 12 months.
- Move to list prices when there are 20 paying customers with churn under 8%, scan-to-paid above 3%, and weekly watch has shipped.
- Annual plans ("2 months free") only after self-serve billing exists.
- Never sell a better model as the main reason to pay; sell reach (logged-in pages) and follow-up (compare, fix prompt, watch).

Initial investment and running costs

Walkthru can launch for well under \$100 in cash. The scarce resource is time, not money.

Item	Cost	When	Notes
Domain	~\$10 to \$20 / year	By 5 Oct	Registrar-dependent; not a quote. Run a quick trademark search on "Walkthru" first
Chrome Web Store developer account	\$5 once	Now	Review can take days; submit by 8 Oct
Google Cloud prepayment (Claude)	₹500 to ₹1,000	Session 6	UPI prepayment; covers the measured Claude comparison (~\$1 to \$2) and early paid runs
API server	\$0 or ~\$5 / month	By 6 Oct	Oracle Always Free VM, or a small VPS
Supabase, Vercel, Resend	\$0	Now	Free tiers; watch quotas
Dodo Payments	\$0 setup	Start verification now	4% + \$0.40 per sale; \$1 per refund; \$30 per dispute; payout fees below \$1,000
Headline ad test (optional)	\$0 to \$100	Oct	Experiment A3 on Reddit ads
OpenRouter credit (optional)	\$10	If report volume grows	Lifts the free model limit from 50 to 1,000 requests a day; needs an international card or crypto
Total to launch	~\$30 to \$150		Estimate, plus founder and co-founder time

Costs that grow with usage

- Claude tokens for paid runs (about 8% of revenue at full use, capped at 25%).
- Dodo fees (about 6% of a \$19 sale).
- Supabase Pro (\$25/month) once the free database or storage fills; Resend and Vercel paid tiers much later.
- Free-model capacity is the first ceiling: about 90 free journeys and 200 scans a day for all users combined. Beyond that, free usage waits for tomorrow or runs on paid capacity funded by revenue.

HOUSEKEEPING BEFORE LAUNCH

Rotate the Supabase database password and secret key (both were shared in a chat on 18 Sep). Turn on two-factor login for Supabase, Google Cloud, Dodo and GitHub. Have the legal pages (privacy, terms, security) reviewed.

Go-to-market plan

Built with the go-to-market founder skill for a \$0 to \$500 monthly budget and a 20 October launch. The short version: give away great scans in public, measure whether the extension kills the funnel, and let every shared report and badge advertise the product.

Launch readiness check

- **Minimum to launch:** Instant Scan on a real domain, the extension on the Chrome Web Store, one public demo report, the new pricing page, and paid passes through Dodo. Plus features can wait.
- **Enough users?** Yes. The AI-builder communities below have hundreds of thousands of members between them.
- **Biggest launch risk:** the install step. If people scan but never install, paid plans have no path. Experiment A2 measures this before launch.

Pre-launch (now to 20 Oct): three experiments, no new code

#	Test	How	Pass if
A1	Will anyone pay once?	Post 30 free Instant Scans of real launches in Indie Hackers, r/SaaS and r/SideProject threads over 10 days, with one line: "Want AI users to test your signup? Founding Launch Pack, \$9." Take payment by Dodo payment link or UPI.	5 paid of the first 100 scans
A2	Does the extension kill the funnel?	20 beta users load the unpacked extension with a 3-line guide; count clicked, installed, finished.	5 or more finish a run within 48 h (else build the cloud runner)
A3	Is GEO the stronger hook?	Same offer, two headlines: "Can ChatGPT read your site?" vs "Find where users get stuck". Count scan starts by link tag.	GEO headline gets 1.5x more starts, then leads the landing page

Audience building posts (before launch):

1. Indie Hackers: "I scanned 30 AI-built apps. 22 are invisible to ChatGPT. Here is why." (data post, links to public reports with owners' permission)
2. X / Twitter build-in-public thread: a real before-and-after rerun on the showcase site, with the fix prompt.
3. r/SideProject: the tool itself, as a free scan, with the story of building it.
4. Lovable and Bolt communities (Discord, r/lovable): "How to make your Lovable app readable by Google and ChatGPT", with the Vite rendering fix.
5. dev.to / Hashnode: "The 7 things AI crawlers need from your site", from our GEO categories.

Waitlist: only for Plus ("weekly watch and MCP, opening in November"). Everything else is usable on day one.

Assets: a 60 to 90 second demo video (paste URL → GEO shock "AI sees 6 words" → extension journey → fix prompt → rerun shows fixed); one public demo report; 3 beta quotes collected from A1 and A2 users with permission.

Launch day (Tue 20 Oct)

Platform	Tactic
----------	--------

Hacker News	Show HN (it is usable without sign-up, which HN rules favour): "Show HN: Walkthru, a launch check that shows what ChatGPT sees on your site". The founder answers every technical comment for the first 3 hours: the safety model, grounding, passive-only security. Never ask for upvotes; HN rules forbid it.
Product Hunt	Launch at 12:01 am Pacific (Tuesday to Thursday are busiest). A hunter is not required. Prepare gallery images from the real report, the demo video, a founding-price offer and a maker comment telling the story. The co-founder runs the day: replies, thank-yous, sharing to communities.
Reddit	r/SaaS limits self-promotion to once every 60 days and has a weekly share thread: use the thread. r/SideProject and r/indiehackers accept projects with a story. Format: problem, what we found scanning 30 apps, free link, no hard sell.
X	Thread: 1) the GEO shock screenshot, 2) how the test user works, 3) safety rules, 4) before/after rerun, 5) link. Tag the builder communities (Lovable, Bolt, Cursor) only where the content genuinely helps their users.

First 90 days: channels ranked

#	Channel	CAC (Estimate)	Time to results	First action this week
1	Free scans posted in founder communities (IH, r/SaaS weekly thread, r/SideProject, Lovable and Bolt communities)	\$0, ~1 h per customer	Days	Scan 10 apps from this week's launches and send owners their report
2	Badge and public-report backlinks		\$0 Weeks to months	Ship V8 and ask every beta user to add the badge
3	Launch platforms (Show HN, Product Hunt)		\$0 One day spike	Draft the Show HN text and PH gallery
4	Search content around GEO ("can ChatGPT read my site", "Lovable SEO")	\$0, writing time	2 to 3 months	Publish one guide with a free-scan call to action
5	Agency outreach (small web agencies, Webflow and Framer freelancers)	\$0 to \$50	1 to 2 months	Send 20 agencies a free white-label-style report of one client site

Content ideas: "What ChatGPT sees on 30 AI-built apps" (data study); "Lovable/Bolt SEO and GEO checklist"; "Why your Supabase signup emails never arrive"; "Security headers in 5 minutes for Vercel apps"; "We let an AI try to sign up to 20 Product Hunt launches". Each becomes an X thread, a Reddit post and a newsletter pitch.

Partnerships: (1) template and course creators in the Lovable, Bolt and Cursor ecosystems, who can bundle a free scan in their launch checklists; (2) MCP directories and the Cursor/Claude Code tool listings once the MCP server ships. **Product-specific tactic:** "Launch Ready" badge leaderboard of the week's launches, with owners' permission.

Metrics for the first 90 days

Metric	Day 30	Day 60	Day 90	Tool	If below target
Instant Scans per week	150	300	500	Supabase runs table	Post more public scans; test the GEO headline
Scan → signup	15%	20%	25%	Supabase auth + runs	Make the locked fix prompt more visible in reports
Install → first finished run	25%	35%	40%	Runs by user	Onboarding video; build the cloud runner
Paying customers	5	15	30	Dodo dashboard	Interview 5 non-payers; revisit the offer

Paid users who rerun within 14 days	40%	50%	60%	Runs with a comparison	Rerun reminder email; weekly watch sooner
-------------------------------------	-----	-----	-----	------------------------	---

All targets are Estimates. Walkthru has no product analytics yet; these numbers come from the database. Adding privacy-safe analytics is a pre-launch task.

Budget (\$0 to \$500 a month)

- **Free:** community posts, launches, content, badge loop, founder-led outreach. This is 90% of the plan.
- **Worth paying for:** the domain (\$15/yr), Chrome Web Store (\$5), Claude comparison (~\$2), and later a \$25/month Supabase upgrade.
- **Best single spend under \$200: \$100 on the A3 Reddit headline test.** It decides the landing page headline, which every later visitor sees, so it compounds.

What we build next

Quality first, then payments, then deploy and launch, then Plus. Dates are from the roadmap; the buffer week exists for Chrome Web Store review and Dodo verification, which are outside our control.

When	Work	What it means
Now to 8 Oct	Polish and harden	Commit the Launch Ready score and badge (V8); production-ready agent state (CHECKPOINTER=postgres); a UI spacing pass; fix the free-model clarity misfire and local http noise; configure PageSpeed.
6 to 8 Oct	Deploy, domain, store	Buy the domain; web on Vercel; API on the Oracle VM with process supervision; production CORS and extension origin; error logging and health checks; submit to the Chrome Web Store on 8 Oct.
6 to 8 Oct	Security tightening	Rotate Supabase secrets; two-factor on every admin account; one API process so the per-address scan limit holds (or move it to the database); verify row-level security on every table; audit extension permissions; legal pages reviewed; Google and GitHub sign-in enabled.
11 to 13 Oct	Dodo payments (V10)	Test mode first: access request → founder approval → private one-use checkout → signed, idempotent webhook inserts the pass → refunds suspend it. Tests prove duplicate, forged and out-of-order webhooks do nothing and a redirect without a webhook grants nothing. Then live mode once Dodo verifies the business.
6 to 15 Oct	Claude on for paid plans	Google Cloud project with UPI prepayment; measured comparison of Claude Haiku 4.5 vs the free chain on the same journeys (\$1 to \$2); only then claim paid reports are better.
14 to 18 Oct	Close-out	Evidence and PDF verified with real screenshots (V7); signup funnel numbers (V18); landing copy review (V19); new landing page and pricing (V9); stranger test in production.
20 Oct	Launch	Show HN, Product Hunt, Indie Hackers, r/SaaS, r/SideProject. 27 Oct at the latest.
21 Oct to 1 Nov	Weekly watch + deploy hook (V12)	A weekly server-side scan per saved site; a Vercel or Netlify deploy hook; email only when something changes.
2 to 8 Nov	Walkthru MCP server (V20)	Plus users paste a URL and a personal API key into Claude Code or Cursor. Tools: scan_site , get_report , get_fix_prompt , rerun , list_runs . Keys shown once, stored hashed, revocable, same limits as the web.
9 to 15 Nov	Competitor side by side (V21), branded PDF, open Plus	Passive scans of up to 3 competitor URLs next to yours; white-label PDFs; Plus waitlist opens by 22 Nov.
16 to 22 Nov	Custom test users (V13)	Define your own persona; several personas merged in one report.
Late Nov (if needed)	Cloud runner (V14)	Only if fewer than 25% of beta users finish a run: public journeys without installing anything.
Dec onward	Self-serve checkout	After 10 reconciled passes and 30 days of cost data. Then: preview-deploy check (GitHub Action), findings to GitHub Issues or Linear, AI citation tracking as a Plus add-on.

Team collaboration on Plus

The pricing research suggested Plus include **3 seats sharing reports**; the current plan table does not list seats. Today each account is one person, and reports are shared by public link. Shared team workspaces (invite teammates, shared sites and runs, roles) would need a new team table and permission rules. **Open decision:** add seats to Plus before it opens in November, or keep Plus single-seat and use MCP plus shared links as the "team" story. Recommendation: ship Plus with shared links and MCP first, add 3 seats only if agencies on the waitlist ask for it.

Risks and open decisions

Risk	Level	Mitigation
Nobody pays	High	A1 before billing code: 5 of 100 scans at \$9. If it fails, change the offer before building more.
Extension install kills the funnel	High	A2 gate at 25%; cloud runner for public journeys.
CanaryUsers or Meerkat add SEO and GEO	Medium	Ship first, own "launch check" and the no-password logged-in story, build the badge loop.
A platform (Vercel, Lovable) builds this in	Medium	Be the neutral check across all builders; partner with their ecosystems.
Free AI capacity runs out	Medium	Daily caps with clear messages; paid runs on paid capacity funded by revenue.
A wrong AI judgement damages trust	Medium	Grounding; tune the clarity rule; Claude for paid; "which models ran" on every report.
The agent does something harmful	Low	Two-layer code rules, verified domains, one confirmed send, passive security only.
Chrome Web Store or Dodo delays	Medium	Submit 8 Oct; start Dodo verification now; one buffer week.
Name or trademark conflict	Unknown	Search before buying the domain.

Decisions still open

- Keep list prices at \$19 and \$49, or adjust after A1.
- Refund policy and legal wording.
- Team seats on Plus.
- Whether Launch Pack and Plus are visible while paid access is request-only.
- The funded AI balance and daily and monthly platform caps.
- Which claims the pricing page may make about Claude, after the measured comparison.

The co-founder role

The product is built. The hardest remaining problem is not code; it is getting founders to try it, pay for it and tell others. Walkthru is looking for a business co-founder to own everything that is not engineering.

What the technical founder owns

- The product, the code, the AI and the safety model.
- Infrastructure, security, payments integration, the Chrome extension.
- Shipping the roadmap in chapter 18.
- Answering technical questions in launches (Show HN).

What the business co-founder owns

- **Customers:** finding, talking to and signing the first 50.
- **Distribution:** communities, launches, content, partnerships.
- **Rollout:** launch day, onboarding emails, the support inbox.
- **Business:** pricing tests, Dodo operations, refunds, legal review, metrics.
- **Voice of the customer:** turning conversations into the next feature.

The first 90 days

Period	Focus	Done when
Weeks 1 to 2	Run experiment A1: post 30 free scans of real launches, collect replies, take the first \$9 payments by payment link. Recruit 20 beta users for A2. Talk to 10 founders about how they check a launch today.	100 scans done, first payments or a clear "no", 10 interview notes
Weeks 3 to 4	Launch prep: Product Hunt page, Show HN draft, demo video script, 3 beta quotes, launch-day community list, onboarding email sequence.	Launch kit ready by 18 Oct
Week 5	Launch week: run Product Hunt and communities; the technical founder handles HN and bugs.	First 10 paying customers
Weeks 6 to 13	Weekly content, agency outreach (20 a week), badge adoption, founder-approved pass operations, churn interviews, Plus waitlist.	30 paying, Plus waitlist of 50

How we would measure the partnership

- Weekly: scans, signups, installs, finished runs, paying customers, revenue, churn (the chapter 17 table).
- A shared Monday review of numbers; one written decision log (the repo already keeps one in `docs/decisions.md`).

Terms to discuss

Equity and vesting are for the two founders to agree, ideally with advice. Common practice among early-stage startups is founder equity vesting over 4 years with a 1-year cliff, so both people earn their share over time. A trial project first (for example, running A1 together over two weeks) is a low-risk way to test the working relationship.

WHAT YOU GET ON DAY ONE

A working product with 171 passing tests, a clear position against named competitors, a pricing model with healthy margins, a launch date, near-zero running cost, and a technical partner who ships daily. What is missing is exactly the part you would own.

[Placeholders for the founder to fill: founder bio and track record, links to past projects, contact details, preferred working hours and location.]

Sources and notes

Project sources (in the repository)

- `SPEC.md` (plans, GEO rules), `ROADMAP.md` (weak points, dates), `ARCHITECTURE.md` (modules), `payment.md` (billing and costs), `CURRENT_STATE.md` (what is done), `docs/decisions.md` (measurements: model bakeoff, Checkpoint B, Claude costs, latency).
- `founder/competitor-matrix.md`, `founder/pricing-strategy.md`, `founder/validate-idea.md`, plus `founder/go-to-market.md` and `founder/product-brief.md` written with this briefing.
- Screenshots: captured on 24 Sep 2026 from the local build with headless Edge. The two reports are real Instant Scans of the fixture sites on ports 8102 and 8103. The extension panel and "stuck" images are design mocks, labelled where used.

External sources

- Competitor pricing: runmeerkat.com/pricing, canaryusers.ai, useswarm.co, uxia.app, test-lab.ai (Octomind), frase.io (Otterly). Checked 24 Sep 2026.
- UserTesting acquisition: [Thoma Bravo press release](#).
- Lovable users and revenue: [Wikipedia: Lovable](#), getpanto.ai Lovable statistics.
- Vibe-coding market size, Cursor and Replit figures: findskill.ai, [Taskade](https://taskade.com), [Hostinger](https://hostinger.com). Industry estimates; treat as directional.
- Show HN rules: [Hacker News FAQ](#); [Syften guide](#).
- r/SaaS self-promotion rule (once every 60 days, weekly share thread): [OneUp](#), [Soar](#). Check each subreddit's current rules before posting.
- Dodo pricing and payouts: dodopayments.com/pricing, [payout structure](#).
- Google Cloud UPI billing: [Google Cloud docs](#). llms.txt evidence: [Otterly study](#).

How to regenerate this PDF

The source is `docs/briefing/walkthru-briefing.html` with images in `docs/briefing/img`. Print it with headless Edge: `msedge --headless=new --no-pdf-header-footer --print-to-pdf=Walkthru-Briefing.pdf docs/briefing/walkthru-briefing.html`.